

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.preprocessing import StandardScaler

# Phase 1: Data Collection (e.g., UGRansome dataset)

pd.set_option("expand_frame_repr", False)
df= pd.read_csv('/kaggle/input/ugransome-dataset/final(2).csv')
df2 = pd.DataFrame(df)
df2.columns = ['Time', 'Protocol', 'Flag', 'Family', 'Clusters', 'SeedAddress', 'ExpAddress']
```

```
Out[1]:
```

	Time	Protocol	Flag	Family	Clusters	SeedAddress	ExpAddress
<b>0</b>	50	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>1</b>	40	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>2</b>	30	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>3</b>	20	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>4</b>	57	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
...	...	...	...	...	...	...	...
<b>149038</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149039</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149040</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149041</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149042</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ

149043 rows × 14 columns

```
In [2]: # Data cleaning
# Renaming the attack "Bonet" to "Botnet"

df2['Threats'] = df2['Threats'].str.replace('Bonet', 'Botnet')

# Print the modified DataFrame
df2
```

Out[2]:

	Time	Protocol	Flag	Family	Clusters	SeedAddress	ExpAddress
<b>0</b>	50	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>1</b>	40	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>2</b>	30	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>3</b>	20	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
<b>4</b>	57	TCP	A	WannaCry	1	1DA11mPS	1BonuSr7
...	...	...	...	...	...	...	...
<b>149038</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149039</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149040</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149041</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ
<b>149042</b>	33	UDP	AP	TowerWeb	3	1AEoiHYZ	1SYSTEMQ

149043 rows × 14 columns

In [3]:

```
# Phase 2: Data Preparation (feature engineering and data transformation)

# --- Drop all duplicate rows --- #

df2 = df2.drop_duplicates()

# --- Remove negative values from time/timestamp feature --- #

df2['Time'] = df2['Time'] + 11

# adding 11 to each value in the 'Time' column of the DataFrame 'df2'.
#In other words, it's performing an element-wise addition operation on all 1
#increasing each value by 11 units. This is often done in data manipulation
#by a fixed amount

# --- Math transformations to reduce skewness --- #

# --- Log transformation applied to column NETFLOW_BYTES --- #
# A log transformation involves taking the natural logarithm (base e) of each
#Logarithmic transformations are often used to reduce the impact of extreme
#closely to a normal distribution. They are particularly useful when dealing
#where the tail of the distribution is elongated on the right side.

#The np.log() function is a common way to perform a logarithmic transformation
#The + 1 added to the data points is often used to avoid issues with taking
#It's a common practice to add a small constant like 1 to the data before applying

#By applying a log transformation to a feature, you're essentially compressing
#which can help in cases where the data exhibits a rightward skew, making it
#or modeling techniques that assume normally distributed data.
```

```

df2['Netflow_Bytes'] = np.log(df2['Netflow_Bytes']+1)

# --- Square root transformation applied to columns USD ---#

#Square Root Transformation: A square root transformation involves taking the
#specified column. In this case, it's applied to the 'USD' column.

#Square root transformations are a type of mathematical transformation used
#Just like logarithmic transformations, square root transformations can help
#a normal distribution.

#The np.sqrt() function is used to calculate the square root.
#By applying a square root transformation to the 'USD' column, the code is made
#and more suitable for certain statistical analyses or modeling techniques that
#require data to be more symmetric. It's a common technique used in data preprocessing
#analysis or modeling

df2['USD'] = np.sqrt(df2['USD'])

# --- Yeo Johnson transformation applied to columns BTC---#

#Yeo-Johnson transformation is being applied to the 'BTC' column in the Data
#This transformation is used to modify the data in the 'BTC' column to make
#The Yeo-Johnson transformation is a mathematical transformation technique used
#It can be applied to both positive and negative values and is more versatile

#The transformation is performed using the stats.yeojohnson() function from
df2['BTC'], _ = stats.yeojohnson(df2['BTC'])

#--PLOTING TRANSFORMED DATA--#

fig, ax = plt.subplots(figsize=(10, 6))

# Plot the transformed 'USD' column
ax.hist(df2['USD'], bins=50, alpha=0.5, color='blue', label='USD (Square Root)')

# Plot the transformed 'BTC' column
ax.hist(df2['BTC'], bins=50, alpha=0.5, color='green', label='BTC (Yeo-Johnson)')

# Plot the transformed 'Netflow_Bytes' column
ax.hist(df2['Netflow_Bytes'], bins=50, alpha=0.5, color='red', label='Netflow_Bytes')

# Add labels and a legend
ax.set_xlabel('Transformed Values')
ax.set_ylabel('Frequency')
ax.set_title('Distribution of Transformed Columns')
ax.legend()

# Show the plot
plt.show()

```

```

# Create a figure and axis for the plot
fig, ax = plt.subplots(figsize=(10, 6))

# Create a StandardScaler instance
# The StandardScaler is a common preprocessing technique used in machine learning
# It is used to standardize or normalize the features of a dataset by scaling
# deviation of 1.

# Standardizing the features is useful because it makes different features more comparable
# that are sensitive to the scale of the input data, such as many machine learning algorithms
# In the code provided, scaler is created as an instance of the StandardScaler
# the specified columns in the df2 DataFrame using the fit_transform method,

scaler = StandardScaler()

# Normalize each column's features
df2_normalized = df2.copy()
df2_normalized[['USD', 'BTC', 'Netflow_Bytes']] = scaler.fit_transform(df2[['USD', 'BTC', 'Netflow_Bytes']])

# Plot the density of the normalized 'USD' column
sns.kdeplot(df2_normalized['USD'], color='blue', label='USD (Square Root)',

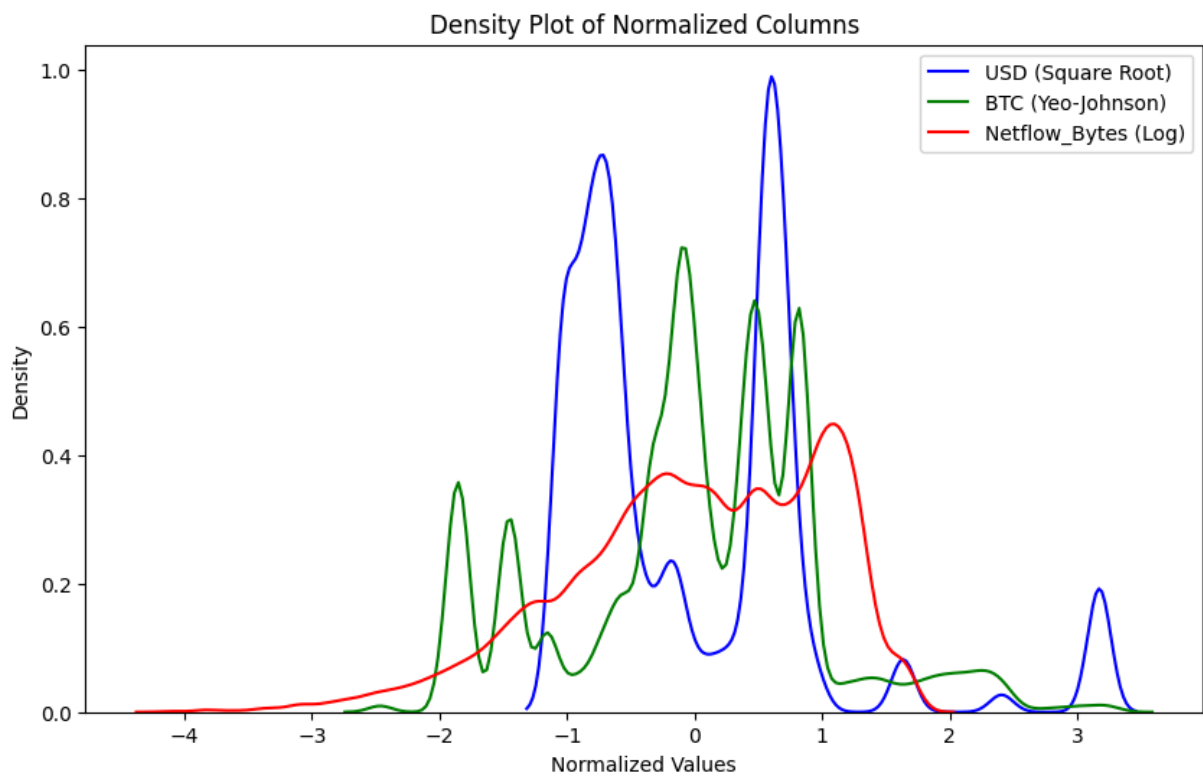

# Plot the density of the normalized 'BTC' column
sns.kdeplot(df2_normalized['BTC'], color='green', label='BTC (Yeo-Johnson)',

# Plot the density of the normalized 'Netflow_Bytes' column
sns.kdeplot(df2_normalized['Netflow_Bytes'], color='red', label='Netflow_Bytes (Log)')

# Add labels and a legend
ax.set_xlabel('Normalized Values')
ax.set_ylabel('Density')
ax.set_title('Density Plot of Normalized Columns')
ax.legend()

# Show the plot
plt.show()

```



```
In [4]: # Phase 3: Data Visualization

# --- Count visualizations --- #

# Categorical count visualizations

# Protocol count
```

Loading [MathJax]/extensions/Safe.js countplot(x=df2['Protocol'], data=df2)

```

plt.title('Bar Graph of Protocol')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_y() + p.get_height() / 2.),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 1),
                textcoords='offset points')

plt.show()

# Flag count

ax = sns.countplot(x=df2['Flag'], data=df2)
plt.title('Bar Graph of Flag')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_y() + p.get_height() / 2.),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 1),
                textcoords='offset points')

plt.show()

# Family count

plt.figure(figsize=(15, 6))
ax = sns.countplot(x=df2['Family'], data=df2)
plt.title('Bar Graph of Family')
plt.xticks(rotation=45)
plt.xticks(fontsize=10)

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_y() + p.get_height() / 2.),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 1),
                textcoords='offset points')

plt.show()

# Clusters count

ax = sns.countplot(x=df2['Clusters'], data=df2)
plt.title('Bar Graph of Clusters')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_y() + p.get_height() / 2.),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 1),
                textcoords='offset points')

plt.show()

# SeedAddress count

ax = sns.countplot(x=df2['SeedAddress'], data=df2)
plt.title('Bar Graph of SeedAddress')
plt.xticks(rotation=45)

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_y() + p.get_height() / 2.),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 1),
                textcoords='offset points')

```

```

        ha='center', va='center', fontsize=10, color='black', xytext=
        textcoords='offset points')

plt.show()

# ExpAddress count

ax = sns.countplot(x=df2['ExpAddress'], data=df2)
plt.title('Bar Graph of ExpAddress')
plt.xticks(rotation=45)

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p
        ha='center', va='center', fontsize=10, color='black', xytext=
        textcoords='offset points')

plt.show()

# IPAddress count

ax = sns.countplot(x=df2['IPAddress'], data=df2)
plt.title('Bar Graph of IPAddress')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p
        ha='center', va='center', fontsize=10, color='black', xytext=
        textcoords='offset points')

plt.show()

# Threats count

ax = sns.countplot(x=df2['Threats'], data=df2)
plt.title('Bar Graph of Threats')
plt.xticks(rotation=45)

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p
        ha='center', va='center', fontsize=10, color='black', xytext=
        textcoords='offset points')

plt.show()

# Port count

ax = sns.countplot(x=df2['Port'], data=df2)
plt.title('Bar Graph of Port')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p
        ha='center', va='center', fontsize=10, color='black', xytext=
        textcoords='offset points')

plt.show()

# Radiation count

```

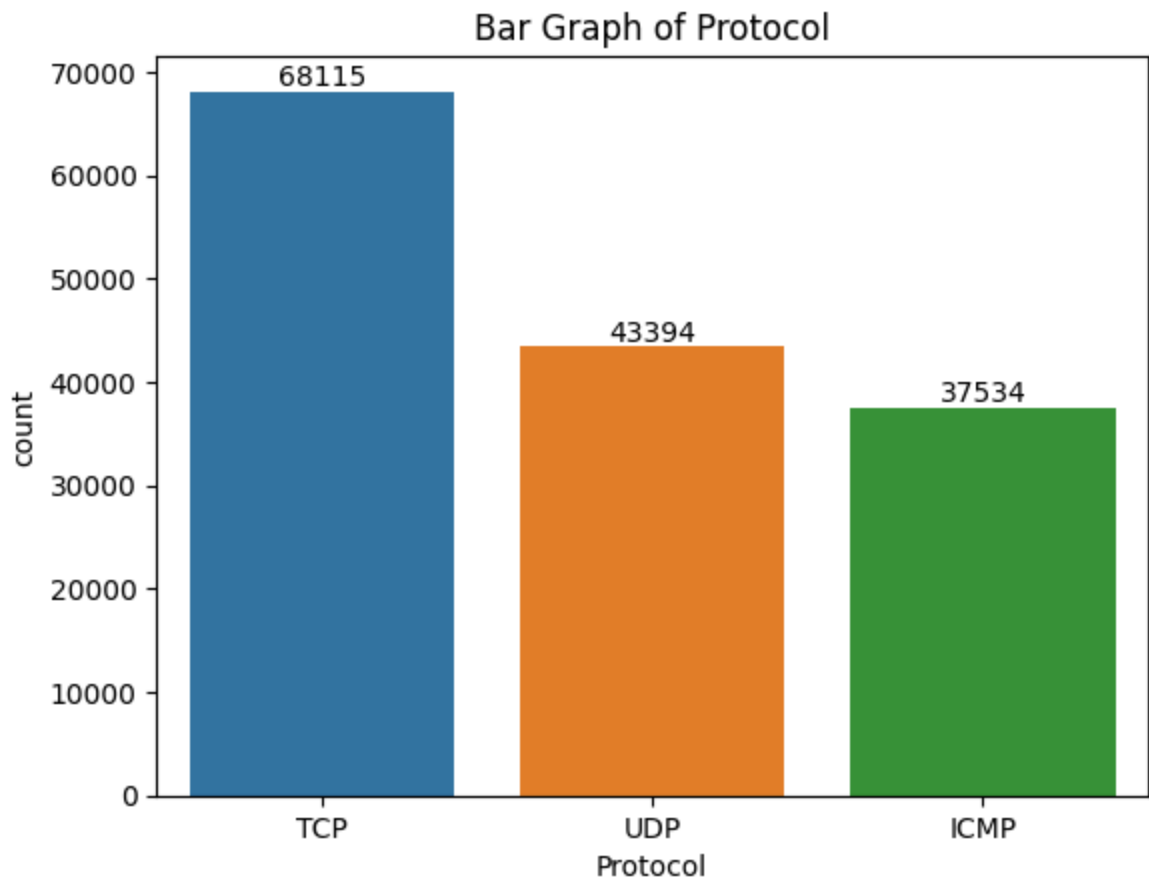
```

ax = sns.countplot(x=df2['Prediction'], data=df2)
plt.title('Bar Graph of Prediction')

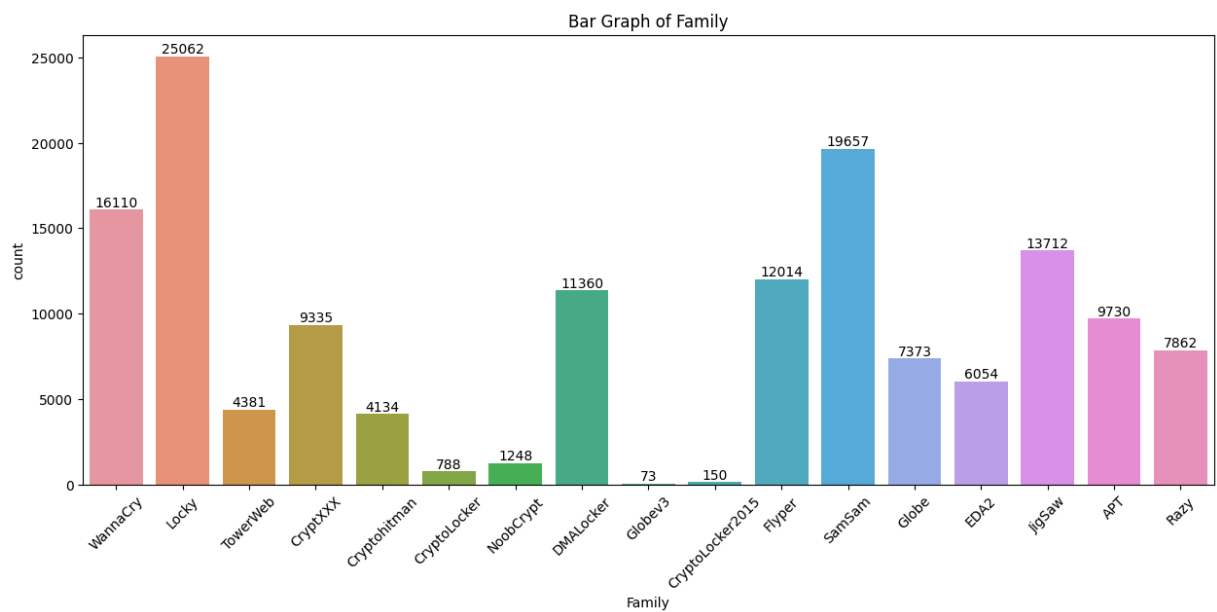
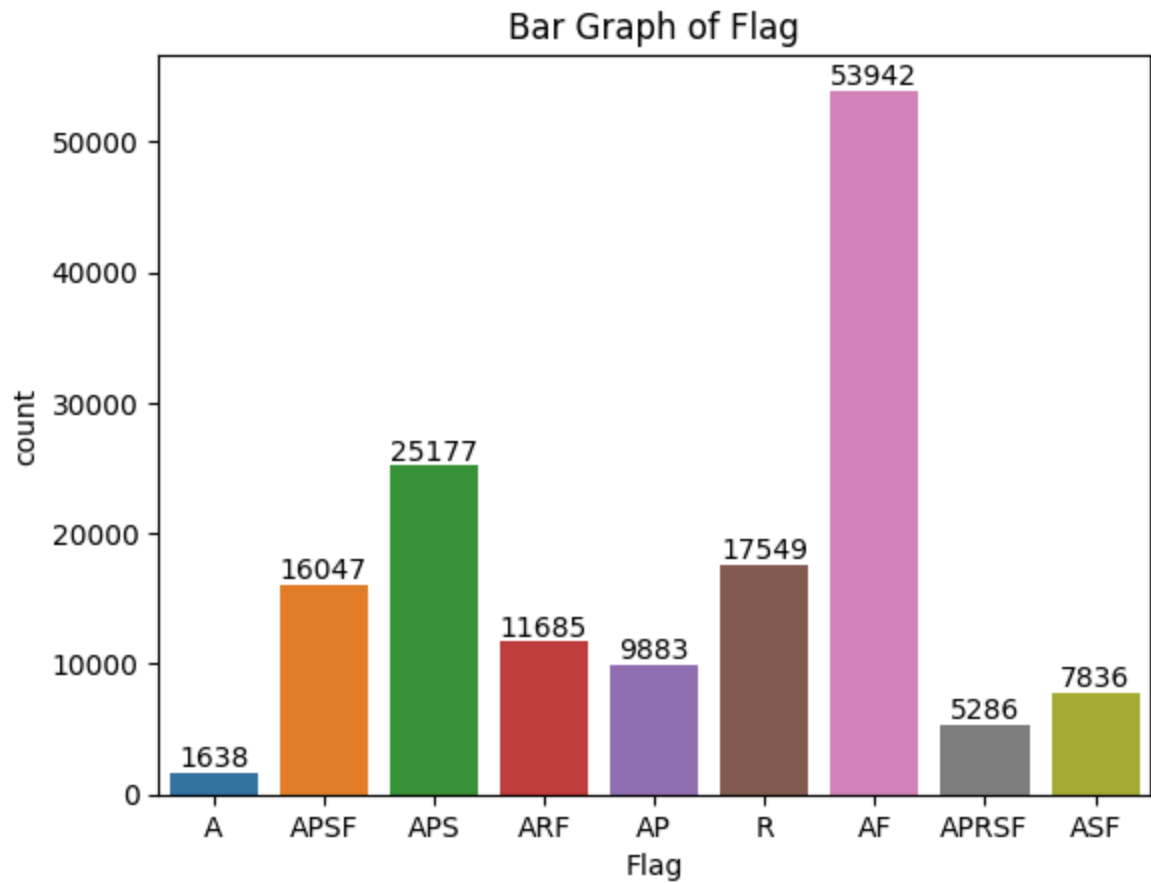
for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='center', fontsize=10, color='black', xytext=(0, 10),
                textcoords='offset points')

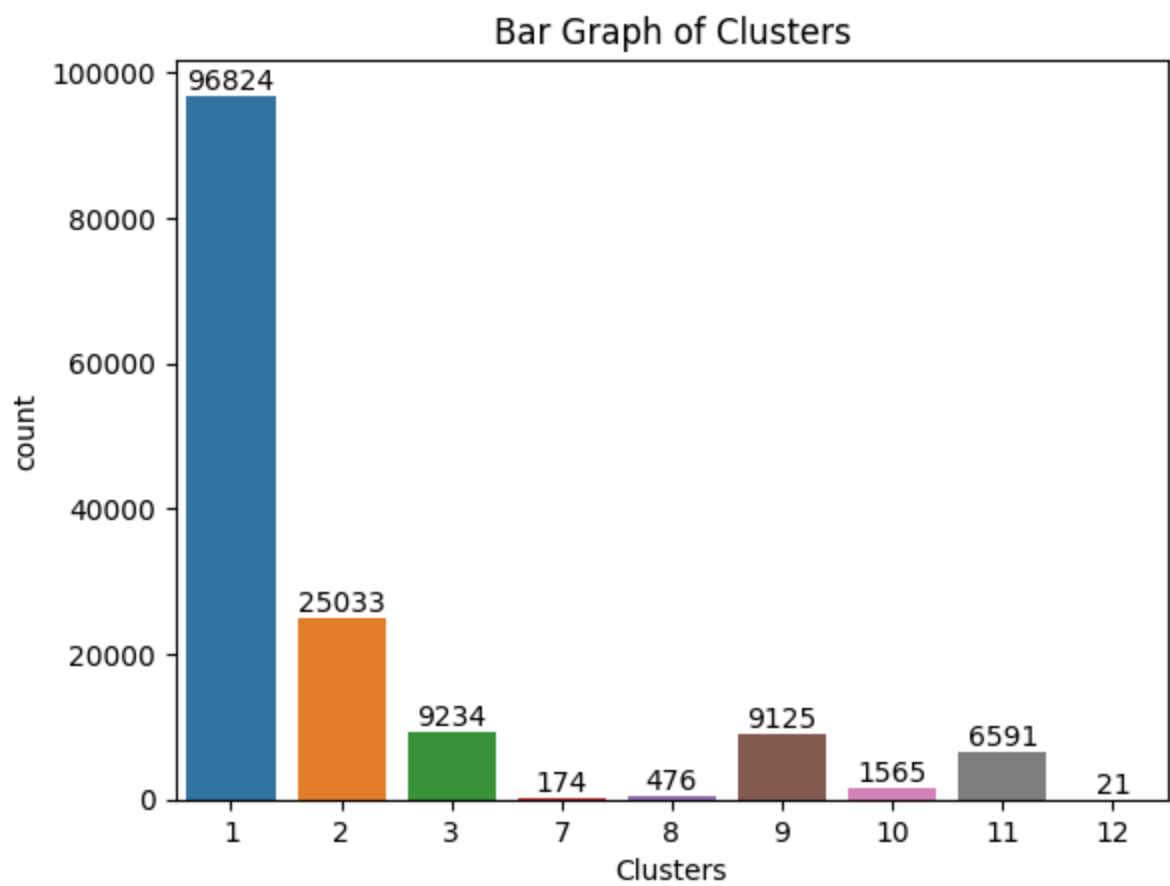
plt.show()

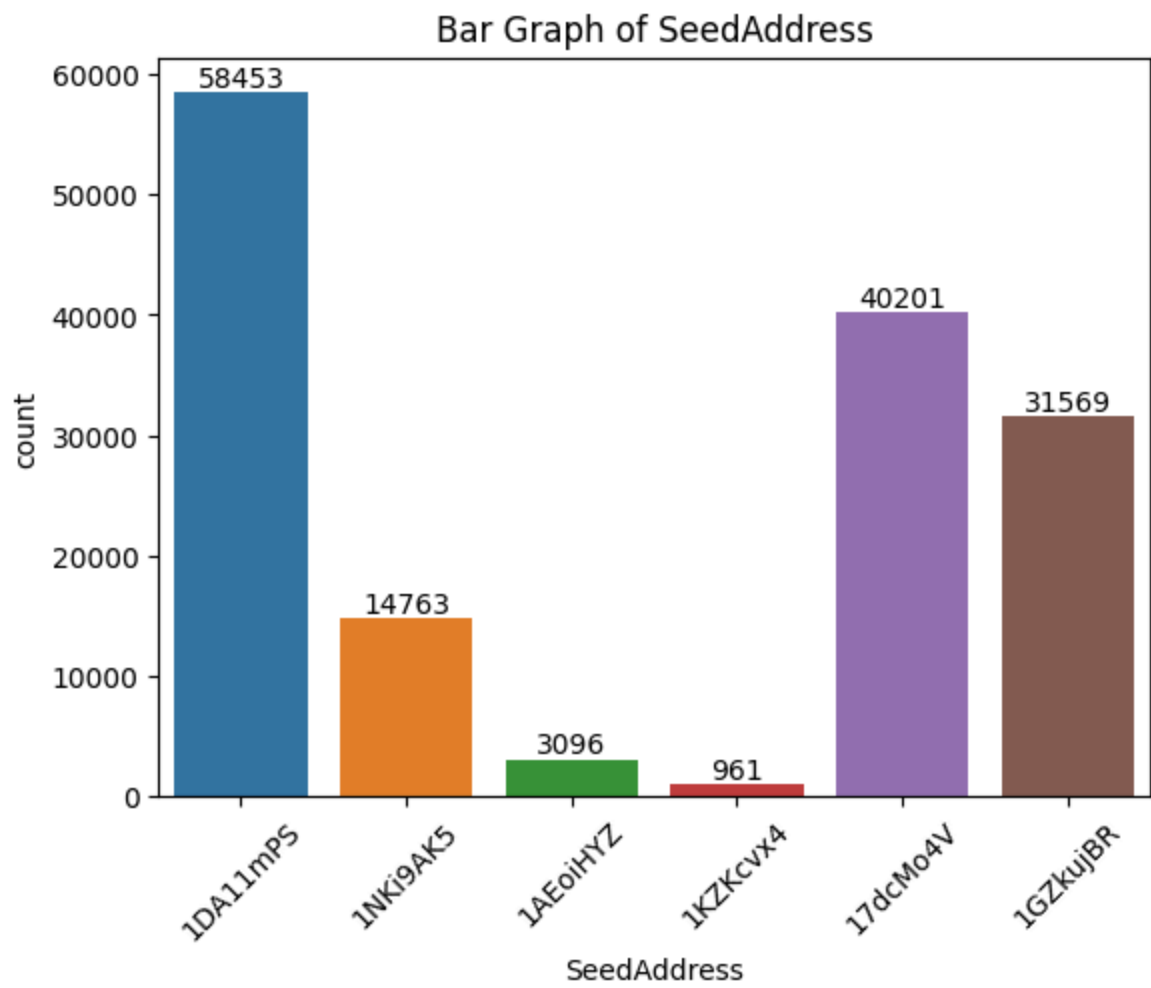
```

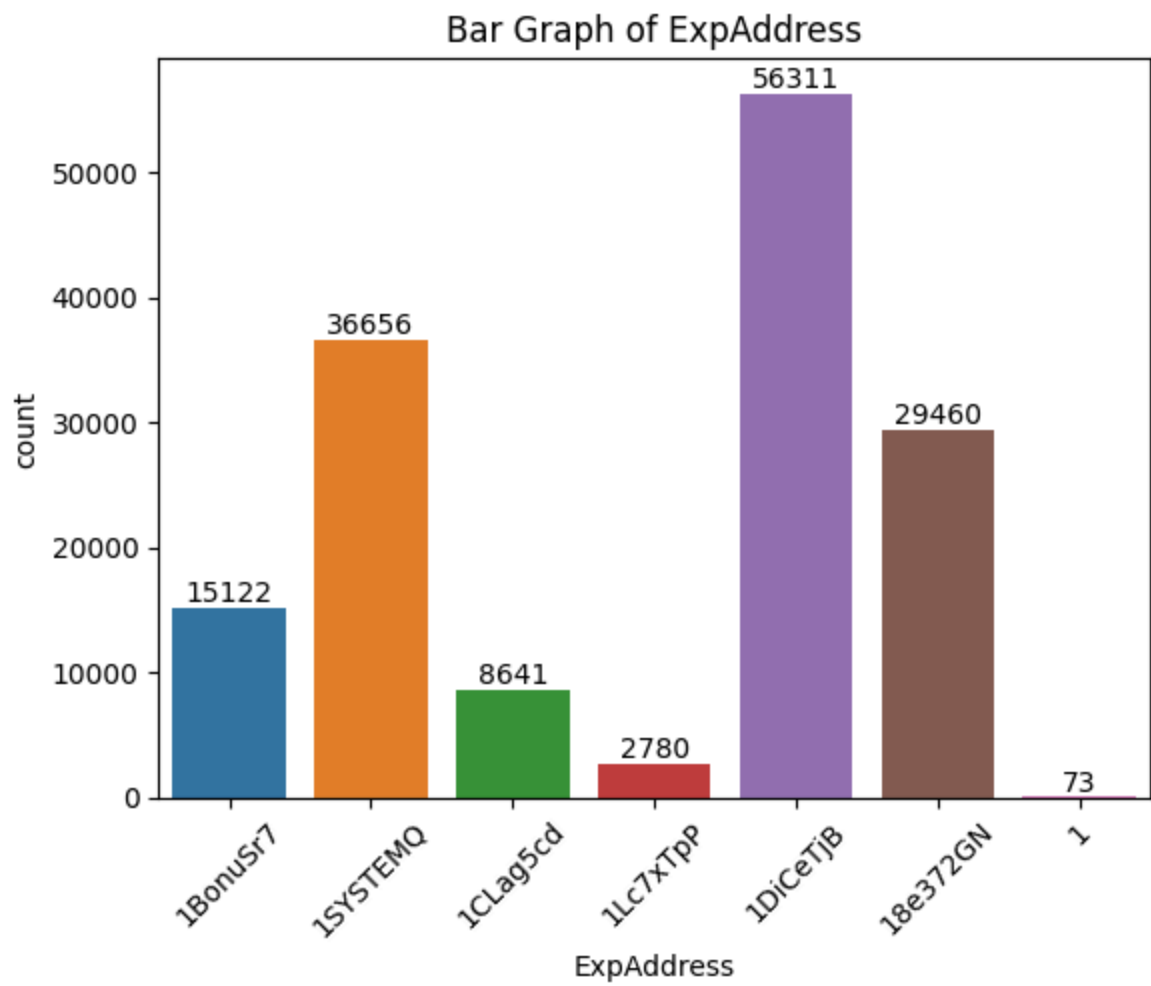





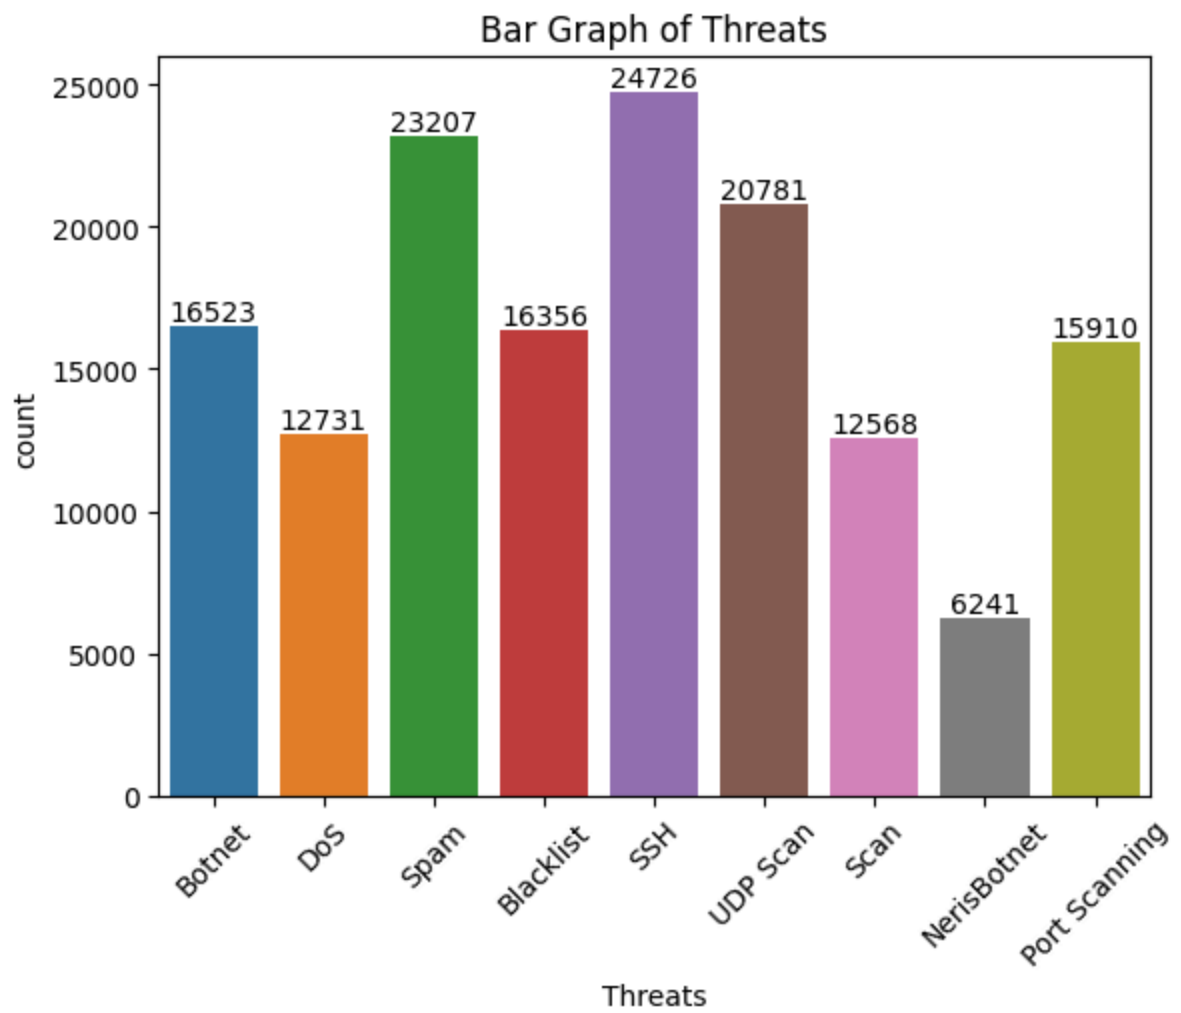



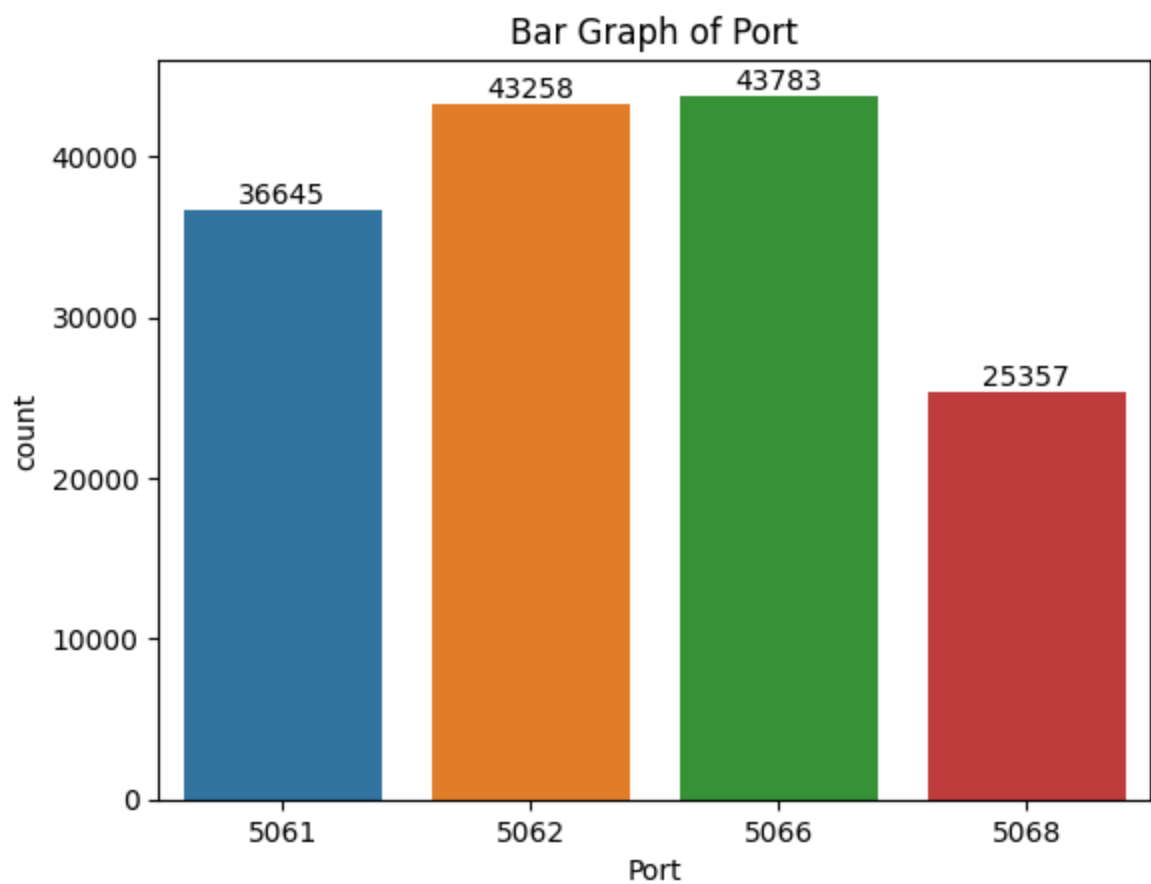












```

In [5]: # --- Numeric visualizations (count, mean and standard deviation) --- #

# Time

feature = 'Time'
data = df2[feature]
mean = np.mean(data)
std_dev = np.std(data)

ax = sns.histplot(data, bins=30, kde=True, color='skyblue', edgecolor='black')
ax.lines[0].set_color('black')

plt.axvline(mean, color='red', linestyle='dashed', linewidth=1, label=f'Mean')
plt.axvline(mean - std_dev, color='green', linestyle='dashed', linewidth=1, label=f'-1σ')
plt.axvline(mean + std_dev, color='orange', linestyle='dashed', linewidth=1, label=f'+1σ')
plt.axvline(std_dev, color='blue', linestyle='dotted', linewidth=1, label=f'σ')

plt.legend(loc='upper right')

plt.title(f'Histogram of {feature}')
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.show()

# # BTC

feature = 'BTC'
data = df2[feature]
mean = np.mean(data)
std_dev = np.std(data)

ax = sns.histplot(data, bins=30, kde=True, color='skyblue', edgecolor='black')
ax.lines[0].set_color('black')

plt.axvline(mean, color='red', linestyle='dashed', linewidth=1, label=f'Mean')
plt.axvline(mean - std_dev, color='green', linestyle='dashed', linewidth=1, label=f'-1σ')
plt.axvline(mean + std_dev, color='orange', linestyle='dashed', linewidth=1, label=f'+1σ')
plt.axvline(std_dev, color='blue', linestyle='dotted', linewidth=1, label=f'σ')

plt.legend(loc='upper right')

plt.title(f'Histogram of {feature}')
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.show()

# # USD

feature = 'USD'
data = df2[feature]
mean = np.mean(data)
std_dev = np.std(data)

ax = sns.histplot(data, bins=30, kde=True, color='skyblue', edgecolor='black')
ax.lines[0].set_color('black')

```



```

plt.axvline(mean, color='red', linestyle='dashed', linewidth=1, label=f'Mean')
plt.axvline(mean - std_dev, color='green', linestyle='dashed', linewidth=1,
plt.axvline(mean + std_dev, color='orange', linestyle='dashed', linewidth=1,
plt.axvline(std_dev, color='blue', linestyle='dotted', linewidth=1, label=f'

plt.legend(loc='upper right')

plt.title(f'Histogram of {feature}')
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.show()

# Netflow_Bytes

feature = 'Netflow_Bytes'
data = df2[feature]
mean = np.mean(data)
std_dev = np.std(data)

ax = sns.histplot(data, bins=30, kde=True, color='skyblue', edgecolor='black'
ax.lines[0].set_color('black')


plt.axvline(mean, color='red', linestyle='dashed', linewidth=1, label=f'Mean')
plt.axvline(mean - std_dev, color='green', linestyle='dashed', linewidth=1,
plt.axvline(mean + std_dev, color='orange', linestyle='dashed', linewidth=1,
plt.axvline(std_dev, color='blue', linestyle='dotted', linewidth=1, label=f'

plt.legend(loc='upper right')

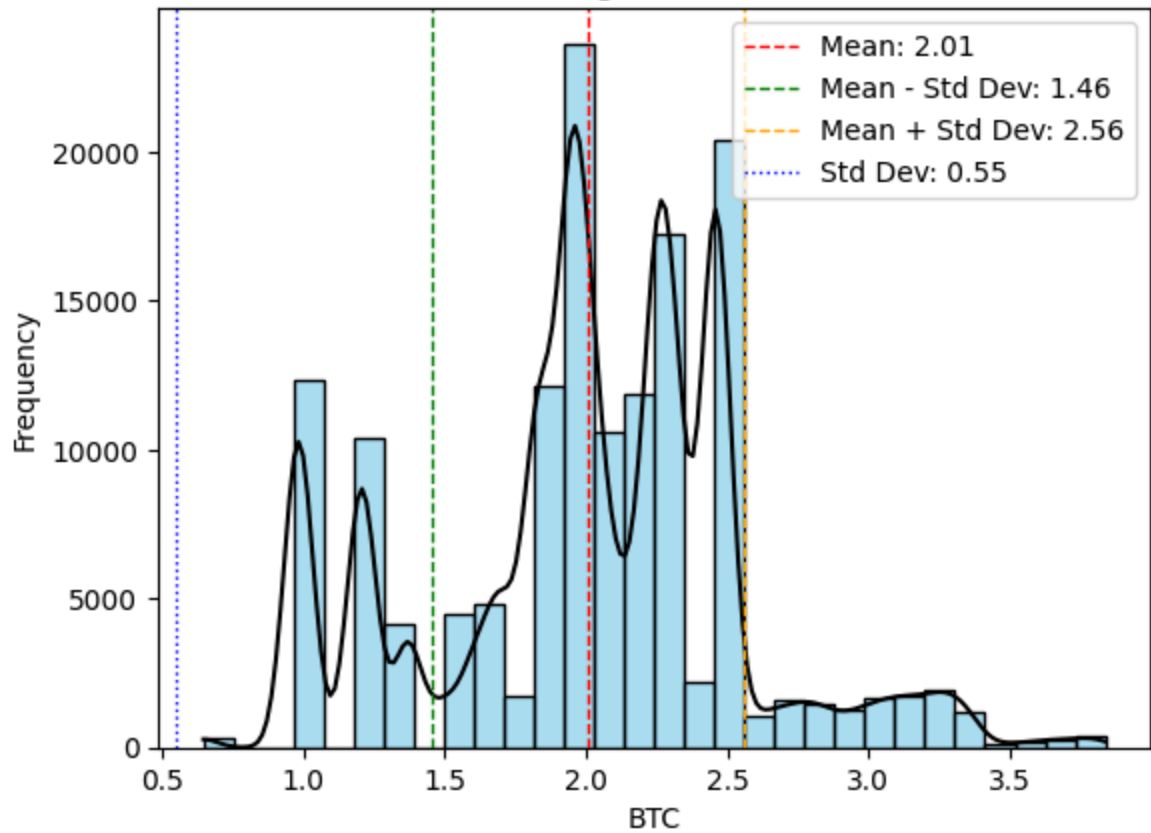
plt.title(f'Histogram of {feature}')
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.show()


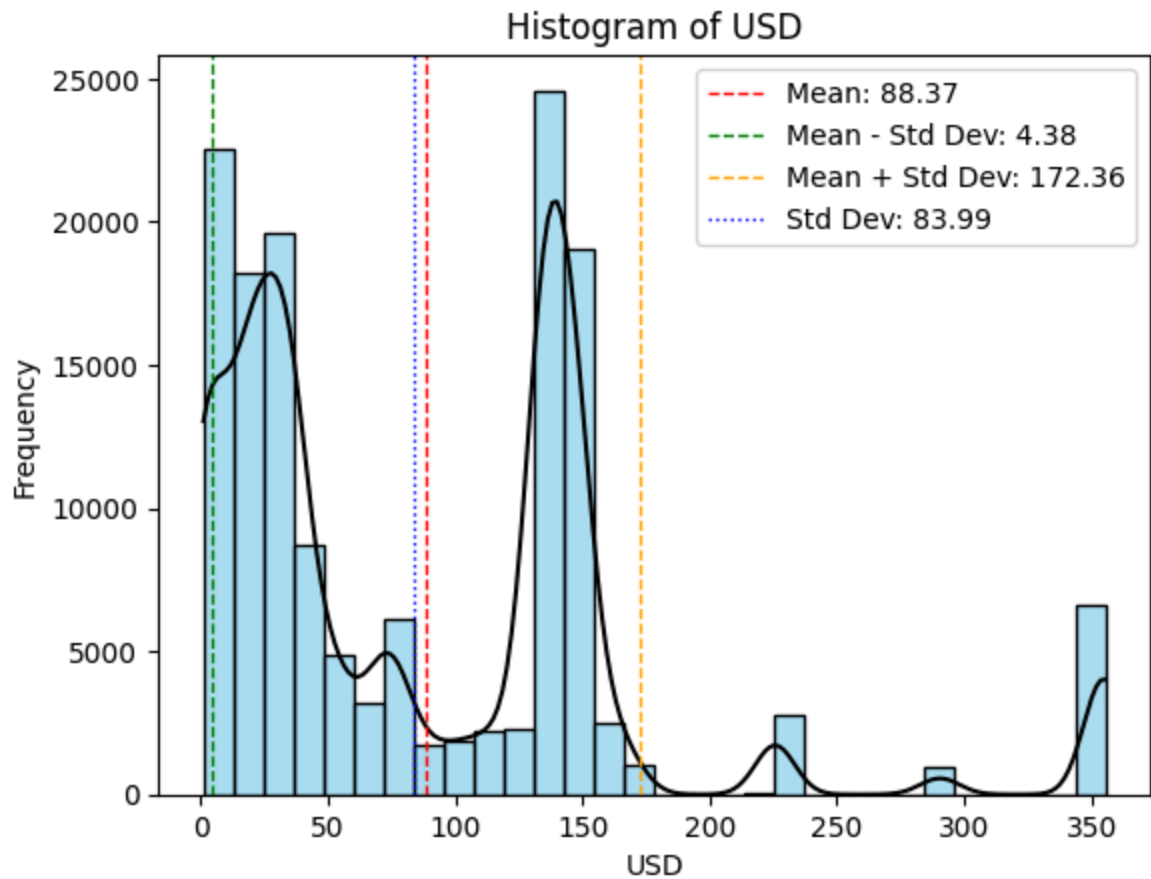
```

### Histogram of Time



### Histogram of BTC





```
In [6]: #The preprocessing module in scikit-learn provides various tools and techniques for
#feeding it into machine learning models.
#This preprocessing is crucial to improve the quality of your data and the performance of the model.

from sklearn import preprocessing
#The code segment uses scikit-learn's LabelEncoder to transform categorical data into numerical labels.
#Each categorical column, such as 'Protocol,' 'Flag,' 'Family,' 'SeedAddress,' 'ExpAddress,' 'IPaddress,'
#'Prediction,' is encoded into unique numeric labels.
#This preprocessing step is essential for machine learning algorithms, as they require numerical input
#instead of categorical labels.

lab_encoder = preprocessing.LabelEncoder() # transform categorical data into numerical labels
df2['Protocol'] = lab_encoder.fit_transform(df2['Protocol'])
df2['Flag'] = lab_encoder.fit_transform(df2['Flag'])
df2['Family'] = lab_encoder.fit_transform(df2['Family'])

df2['SeedAddress'] = lab_encoder.fit_transform(df2['SeedAddress'])
df2['ExpAddress'] = lab_encoder.fit_transform(df2['ExpAddress'])
df2['IPaddress'] = lab_encoder.fit_transform(df2['IPaddress'])
df2['Threats'] = lab_encoder.fit_transform(df2['Threats'])
df2['Prediction'] = lab_encoder.fit_transform(df2['Prediction'])
df2
```

```
Out[6]:
```

	Time	Protocol	Flag	Family	Clusters	SeedAddress	ExpAddress	
<b>0</b>	61	1	0	16	1	2	2	0.0
<b>1</b>	51	1	0	16	1	2	2	0.0
<b>2</b>	41	1	0	16	1	2	2	0.0
<b>3</b>	31	1	0	16	1	2	2	0.0
<b>4</b>	68	1	0	16	1	2	2	0.0
...	...	...	...	...	...	...	...	...
<b>149038</b>	44	2	2	15	3	1	6	3.0
<b>149039</b>	44	2	2	15	3	1	6	3.0
<b>149040</b>	44	2	2	15	3	1	6	3.0
<b>149041</b>	44	2	2	15	3	1	6	3.0
<b>149042</b>	44	2	2	15	3	1	6	3.0

149043 rows × 14 columns

```
In [7]: #The train_test_split function from scikit-learn is used to split a dataset into
#a training set and a testing (or validation) set. This function is commonly used
#to assess the performance of a model on unseen data. It takes as input the dataset
#and labels (y), and divides it into training data (X_train and y_train) used for
#testing data (X_test and y_test) used to evaluate the model's performance.

from sklearn.model_selection import train_test_split # library for machine learning
```

```

#common procedure in machine learning for splitting a dataset into training
#from scikit-learn. Here's a breakdown of what each line of code does:

X = df2.iloc[:, :-1] #This line selects all rows and all columns of the Data
#It's assuming that the last column contains the target variable or labels,

y = df2.iloc[:, -1] # This line selects all rows but only the last column
#This is to isolate the target variable or labels, and y will contain these

X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.8,
#This line uses the train_test_split function to split the data into training
#Here's a breakdown of the parameters:

#X and y: The feature matrix and target variable.
#train_size=0.8: This parameter specifies that 80% of the data should be used
#(you can adjust this percentage as needed).
#random_state=42: This parameter sets the random seed for reproducibility,
#run the code.

#After running this code, you will have:

#X_train: The feature matrix for training.
#X_test: The feature matrix for testing.
#y_train: The target variable for training.
#y_test: The target variable for testing.
#These subsets can then be used for training and evaluating your machine learning model.

```

```

In [8]: X_train
        X_test
        y_train
        y_test

```

```

Out[8]: 42916      1
        45544      2
        137525     0
        108170      1
        85804      2
        ..
        91256      1
        132188      1
        94999      2
        3431       0
        147946     0
        Name: Prediction, Length: 29809, dtype: int64

```

```

In [9]: #The %%time command is typically used in Jupyter Notebook environments, such as
#It is called a "magic command" and is used to measure the execution time of a code cell
#When you include %%time at the beginning of a cell, it tells Jupyter to measure the execution time
#that cell
%%time

```

```

# Imported models: ensemble, random forest, SVM, Naive Bayes, genetic algorithm
# Imported evaluation metrics: accuracy, precision, recall, f1 score

from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split

from sklearn.ensemble import StackingClassifier #ensmbl method of stacking classifiers
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

from sklearn.tree import DecisionTreeClassifier #estimator in GA
import numpy as np

import warnings
warnings.filterwarnings('ignore')

```

```

In [10]: rf = RandomForestClassifier(n_estimators=100, random_state=42) # It specifies the number of trees in the forest
#In this case, there are 100 trees in the forest

# random_state: This parameter is used to set the random seed for reproducibility
#By setting it to 42, the randomization process will be the same each time it is run
#ensuring consistent results for the Random Forest model.

rf.fit(X_train, y_train)

rf_pred=rf.predict(X_test)

#This code snippet uses the trained Random Forest classifier (rf) to make predictions on the test set
#The predict method takes the test features in X_test as input and produces predictions
#The predictions are stored in the rf_pred variable, which can be used for further analysis
#the model performs on unseen data.

rf_accuracy = accuracy_score(rf_pred, y_test)
rf_report = classification_report(rf_pred, y_test)
rf_matrix = confusion_matrix(rf_pred, y_test)
print('Accuracy of Random Forest : ', round(rf_accuracy, 3))
print('Classification report of Random Forest : \n', rf_report)
print('Confusion Matrix of Random Forest : \n', rf_matrix)

#The accuracy_score function from scikit-learn is used to calculate the accuracy of the model
#compared to the actual labels (y_test). This score measures the proportion of correct predictions
#classification_report: The classification_report function generates a comprehensive report of model
#F1-score, and support for each class in the classification problem. It provides a detailed
#performance for different classes.
#confusion_matrix: The confusion_matrix function computes a confusion matrix, which is a table
#false positive, and false negative counts for the classification results. It helps in understanding
ng and where it might be making errors.

```

*#Finally, the code prints out the accuracy, classification report, and confusion matrix  
#allowing you to evaluate its performance on the test data.*

Accuracy of Random Forest : 0.994

Classification report of Random Forest :

	precision	recall	f1-score	support
0	0.99	0.99	0.99	8400
1	0.99	0.99	0.99	13359
2	1.00	0.99	0.99	8050
accuracy			0.99	29809
macro avg	0.99	0.99	0.99	29809
weighted avg	0.99	0.99	0.99	29809

Confusion Matrix of Random Forest :

```
[[ 8346   49    5]
 [   49 13280   30]
 [    8    44 7998]]
```

```
In [11]: svr = LinearSVC()
svr.fit(X_train, y_train)
svr_pred = svr.predict(X_test)

#a Support Vector Machine (SVM) classifier with a linear kernel (LinearSVC)

#svr = LinearSVC(): An instance of the LinearSVC classifier is created.

#svr.fit(X_train, y_train): The LinearSVC classifier is trained on the training data.
#This step involves finding the hyperplane that best separates the data points
#between them.

#svr_pred = svr.predict(X_test): The trained SVM classifier is used to make predictions.
#These predictions are stored in the svr_pred variable.

svr_accuracy = accuracy_score(svr_pred, y_test)
svr_report = classification_report(svr_pred, y_test)
svr_matrix = confusion_matrix(svr_pred, y_test)
print('Accuracy of SVM : ', round(svr_accuracy, 3))
print('Classification report of SVM : \n', svr_report)
print('Confusion Matrix of SVM : \n', svr_matrix)

#svr_accuracy = accuracy_score(svr_pred, y_test): The accuracy of the SVM classifier
#is calculated by comparing them to the true labels (y_test). The result is
```

```

#svr_report = classification_report(svr_pred, y_test): The classification_report
#classification report, including metrics such as precision, recall, F1-score
#This report is stored in the svr_report variable.

#svr_matrix = confusion_matrix(svr_pred, y_test): The confusion matrix is computed
#true labels (y_test). The confusion matrix provides information about the number of
#false positive, and false negative predictions. It is stored in the svr_matrix

#Finally, the results are printed using print statements:

#The accuracy of the SVM classifier is printed with a rounded value.
#The classification report, which includes precision, recall, F1-score, and
#The confusion matrix, which shows the distribution of true and false predictions
#These metrics help evaluate the performance of the SVM classifier in terms

```

Accuracy of SVM : 0.574

Classification report of SVM :

	precision	recall	f1-score	support
0	0.03	0.59	0.05	369
1	0.68	0.93	0.79	9760
2	0.97	0.39	0.56	19680
accuracy			0.57	29809
macro avg	0.56	0.64	0.47	29809
weighted avg	0.86	0.57	0.63	29809

Confusion Matrix of SVM :

```

[[ 219   73   77]
 [ 442 9122  196]
 [7742 4178 7760]]

```

In [12]: *#Naive Bayes Algorithm*

```

nb = GaussianNB()
nb.fit(X_train, y_train)
nb_pred = nb.predict(X_test)

nb_accuracy = accuracy_score(nb_pred, y_test)
nb_report = classification_report(nb_pred, y_test)
nb_matrix = confusion_matrix(nb_pred, y_test)
print('Accuracy of Naive Bayes : ', round(nb_accuracy, 3))
print('Classification report of Naive Bayes : \n', nb_report)
print('Confusion Matrix of Naive Bayes : \n', nb_matrix)

# Assuming you already have nb_pred and y_test defined

nb_accuracy = accuracy_score(nb_pred, y_test)
nb_report = classification_report(nb_pred, y_test)
nb_matrix = confusion_matrix(nb_pred, y_test)

```



```

print('Accuracy of Naive Bayes : ', round(nb_accuracy, 3))
print('Classification report of Naive Bayes : \n', nb_report)
print('Confusion Matrix of Naive Bayes :\n', nb_matrix)

# Plot the confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.set(font_scale=1.2) # Adjust the font size for better readability
sns.heatmap(nb_matrix, annot=True, fmt="d", cmap="Blues", cbar=False,
            xticklabels=["0:A", "1:S", "2:SS"], yticklabels=["0:A", "1:S", "2:SS"])
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix Heatmap")
plt.show()

```

```

Accuracy of Naive Bayes : 0.777
Classification report of Naive Bayes :
              precision    recall  f1-score   support

     0       0.64       0.67       0.66       7993
     1       0.86       0.88       0.87      13191
     2       0.78       0.72       0.75       8625

 accuracy          0.78       29809
 macro avg       0.76       0.76       0.76       29809
weighted avg       0.78       0.78       0.78       29809

```

```

Confusion Matrix of Naive Bayes :
[[ 5380  1357  1256]
 [ 1109 11543   539]
 [ 1914   473  6238]]
Accuracy of Naive Bayes : 0.777
Classification report of Naive Bayes :
              precision    recall  f1-score   support

     0       0.64       0.67       0.66       7993
     1       0.86       0.88       0.87      13191
     2       0.78       0.72       0.75       8625


 accuracy          0.78       29809
 macro avg       0.76       0.76       0.76       29809
weighted avg       0.78       0.78       0.78       29809

```

```

Confusion Matrix of Naive Bayes :
[[ 5380  1357  1256]
 [ 1109 11543   539]
 [ 1914   473  6238]]

```



```
In [13]: # Assuming you already have nb_pred and y_test defined

nb_accuracy = accuracy_score(nb_pred, y_test)
nb_report = classification_report(nb_pred, y_test, output_dict=True) # Use
nb_matrix = confusion_matrix(nb_pred, y_test)

# Extract support for all classes
labels = [str(label) for label in np.unique(np.concatenate((nb_pred, y_test)))]
support = [nb_report[label]['support'] if label in nb_report else 0 for label in labels]

print('Accuracy of Naive Bayes : ', round(nb_accuracy, 3))
print('Classification report of Naive Bayes : \n', classification_report(nb_pred, y_test))
print('Confusion Matrix of Naive Bayes : \n', nb_matrix)

# Plot support
plt.figure(figsize=(10, 6))
plt.bar(labels, support, width=0.2, label='Support', align='center')

plt.xlabel('Class')
plt.ylabel('Number of Features')
plt.xticks(labels)
plt.legend()
plt.title('Support for Each Class')
plt.show()
```


Accuracy of Naive Bayes : 0.777

Classification report of Naive Bayes :

	precision	recall	f1-score	support
0	0.64	0.67	0.66	7993
1	0.86	0.88	0.87	13191
2	0.78	0.72	0.75	8625
accuracy			0.78	29809
macro avg	0.76	0.76	0.76	29809
weighted avg	0.78	0.78	0.78	29809

Confusion Matrix of Naive Bayes :

```
[[ 5380 1357 1256]
 [ 1109 11543  539]
 [ 1914  473 6238]]
```



In [14]: *# Assuming you already have nb\_pred and y\_test defined*

```
nb_accuracy = accuracy_score(nb_pred, y_test)
nb_report = classification_report(nb_pred, y_test, output_dict=True) # Use
nb_matrix = confusion_matrix(nb_pred, y_test)

# Extract precision and recall for all classes
labels = [str(label) for label in np.unique(np.concatenate((nb_pred, y_test))
precision = [nb_report[label]['precision'] if label in nb_report else 0.0 for label in labels]
recall = [nb_report[label]['recall'] if label in nb_report else 0.0 for label in labels]

print('Accuracy of Naive Bayes : ', round(nb_accuracy, 3))
print('Classification report of Naive Bayes : \n', classification_report(nb_pred, y_test))

# Plot precision and recall
plt.figure(figsize=(10, 6))
```


```
plt.bar(labels, precision, width=0.2, label='Precision', align='center')
plt.bar(labels, recall, width=0.2, label='Recall', align='edge')

plt.xlabel('Class')
plt.ylabel('Score')
plt.xticks(labels)
plt.legend()
plt.title('Precision and Recall for Each Class')
plt.show()
```

Accuracy of Naive Bayes : 0.777

Classification report of Naive Bayes :

	precision	recall	f1-score	support
0	0.64	0.67	0.66	7993
1	0.86	0.88	0.87	13191
2	0.78	0.72	0.75	8625
accuracy			0.78	29809
macro avg	0.76	0.76	0.76	29809
weighted avg	0.78	0.78	0.78	29809



```
In [15]: estimators = [
            ('rf', RandomForestClassifier(n_estimators=1000, random_state=42)),
            ('svr', LinearSVC(random_state=42))
          ]
clf = StackingClassifier(
    estimators=estimators, final_estimator=GaussianNB())
```

```

#Stacking is an ensemble learning method that combines multiple base estimat

#Here's a breakdown of the code:

#estimators: This is a list of tuples, where each tuple contains the name of
#Two base estimators are defined:

#'rf': A Random Forest Classifier with 1000 estimators and a random seed of
#'svr': A Linear Support Vector Classifier (LinearSVC) with a random seed of
#clf: The Stacking Classifier is created using the StackingClassifier class.

#estimators: This parameter receives the list of base estimators defined ear
#final_estimator: This parameter specifies the meta-estimator that combines
#In this case, a Gaussian Naive Bayes (GaussianNB) classifier is used as the
#The Stacking Classifier combines the predictions of the base classifiers (
#This ensemble method can often improve classification performance by levera

```

```

In [16]: clf.fit(X_train, y_train)
pred = clf.predict(X_test)
accuracy = accuracy_score(pred, y_test)

#In this code snippet, the Stacking Classifier (clf) is trained on the train
#After training, the classifier is used to make predictions on the test data
#Finally, the accuracy of the predictions is calculated using scikit-learn's

#The code essentially performs the following steps:

#Trains the Stacking Classifier (clf) using the training data.
#Uses the trained classifier to predict the target labels for the test data.
#Calculates the accuracy of the predictions by comparing them to the true la
#The accuracy variable will contain the accuracy score of the Stacking Class
#This score measures how well the classifier performed in terms of correctly

```

```

In [17]: eb_accuracy = accuracy_score(pred, y_test)
eb_matrix = confusion_matrix(pred, y_test)
eb_report = classification_report(pred, y_test)

print('Accuracy of Ensemble Model : ', round(eb_accuracy, 3))
print('Confusion Matrix of Ensemble Model : ', eb_matrix)
print('Classification Report of Ensemble Model : ', eb_report)

#In this code snippet, the accuracy, confusion matrix, and classification re
#and printed.

#Here's what each part of the code does:

#eb_accuracy: Calculates the accuracy of the ensemble model's predictions by
#using the accuracy_score function.
#eb_matrix: Computes the confusion matrix for the ensemble model's predictio
#The confusion matrix provides information about the true positives, true ne
#eb_report: Generates a classification report for the ensemble model's predi

```

```
#The classification report includes metrics such as precision, recall, F1-score
#Finally, the code prints out the accuracy, confusion matrix, and classification report
#These metrics provide insights into the model's performance in terms of classification
#classify different classes.
```

```
# Assuming you already have pred and y_test defined for your Ensemble Model
```

```
eb_accuracy = accuracy_score(pred, y_test)
eb_matrix = confusion_matrix(pred, y_test)
eb_report = classification_report(pred, y_test)
```

```
print('Accuracy of Ensemble Model : ', round(eb_accuracy, 3))
print('Confusion Matrix of Ensemble Model : \n', eb_matrix)
print('Classification Report of Ensemble Model :\n', eb_report)
```

```
# Plot the confusion matrix as a heatmap
```

```
plt.figure(figsize=(8, 6))
sns.set(font_scale=1.2) # Adjust the font size for better readability
sns.heatmap(eb_matrix, annot=True, fmt="d", cmap="Blues", cbar=False,
            xticklabels=["0:A", "1:S", "2:SS"], yticklabels=["0:A", "1:S", "2:SS"])
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix Heatmap for Ensemble Model")
plt.show()
```

```
Accuracy of Ensemble Model : 0.994
```

```
Confusion Matrix of Ensemble Model : [[ 8352   49    6]
```

```
 [  45 13282   32]
```

```
 [    6    42 7995]]
```

```
Classification Report of Ensemble Model :
```

```
precision    recall  f1-score   support
```

```
0          0.99      0.99      0.99      8407
```

```
1          0.99      0.99      0.99     13359
```

```
2          1.00      0.99      0.99      8043
```

```
accuracy          0.99      0.99      0.99     29809
```

```
macro avg         0.99      0.99      0.99     29809
```

```
weighted avg      0.99      0.99      0.99     29809
```

```
Accuracy of Ensemble Model : 0.994
```

```
Confusion Matrix of Ensemble Model :
```

```
[[ 8352   49    6]
```

```
 [  45 13282   32]
```

```
 [    6    42 7995]]
```

```
Classification Report of Ensemble Model :
```

```
precision    recall  f1-score   support
```

```
0          0.99      0.99      0.99      8407
```


```
1          0.99      0.99      0.99     13359
```

```
2          1.00      0.99      0.99      8043
```

```
accuracy          0.99      0.99      0.99     29809
```

```
macro avg         0.99      0.99      0.99     29809
```

```
weighted avg      0.99      0.99      0.99     29809
```



```
In [18]: #Plot the evaluation metrics of each model in one figure

# Model names
models = ['Random Forest', 'SVM', 'Naive Bayes', 'Ensemble Learning']

# Precision scores
precision = [100, 82, 64, 99]

# Recall scores
recall = [100, 47, 66, 99]

# F1-score scores
f1_score = [100, 65, 77, 99]

# X-axis values (models)
x = range(len(models))

# Create a figure and axis for the plot
fig, ax = plt.subplots(figsize=(10, 6))

# Plot precision scores
ax.plot(x, precision, marker='o', linestyle='--', color='b', label='Precision')

# Plot recall scores
ax.plot(x, recall, marker='o', linestyle='--', color='g', label='Recall')
```

```

# Plot F1-score scores
ax.plot(x, f1_score, marker='o', linestyle='--', color='r', label='F1-Score')

# Set x-axis ticks and labels
ax.set_xticks(x)
ax.set_xticklabels(models, rotation=45)
ax.set_xlabel('Machine Learning Models')


# Set y-axis label
ax.set_ylabel('Scores (%)')

# Set plot title
ax.set_title('Fluctuation of Precision, Recall, and F1-Score for Different M

# Add a legend
ax.legend()

# Show the plot
plt.tight_layout()
plt.grid(True)
plt.show()

```



```

In [19]: # Define the algorithms and their corresponding accuracies
algorithms = ['Random Forest', 'SVM', 'NB', 'Ensemble Learning']
accuracies = [100, 64, 77, 99]

# Create a bar graph
plt.figure(figsize=(10, 6))
plt.bar(algorithms, accuracies, color=['blue', 'red', 'green', 'purple'])
plt.ylim(0, 110) # Set the y-axis limit for better visualization
plt.xlabel('Algorithms')
plt.ylabel('Accuracy (%)')
plt.title('Accuracy of Different Machine Learning Algorithms')
plt.show()

```




```

# Display the accuracy values on top of the bars
for i, v in enumerate(accuracies):
    plt.text(i, v + 2, str(v) + '%', ha='center', va='bottom', fontsize=12)

# Show the graph
plt.tight_layout()
plt.show()

```



In [ ]: